

An Analysis of Multi-armed bandit

Multi-armed bandit

Multi-armed bandit -- Part 1

Hence, the regret $\rho_{\pi}(T)$ for policy π is computed as the difference between $D(T)$ and the cumulative expected reward at step T for policy π :

Multi-armed bandit -- Part 2

Add noise to each arm and pull the one with the highest value 3. Update value: $R_{I(t)}(t+1) = R_{I(t)}(t) + x_{I(t)}(t)$

Multi-armed bandit -- Part 3

There are two predominant settings in BAI: Fixed budget setting: Given a time horizon $T \geq 1$, the objective is to identify the arm with the highest expected reward $a^* \in \arg \max_k \mu_k$ minimizing probability of error δ . Fixed confidence setting: Given a confidence level $\delta \in (0, 1)$, the objective is to identify the arm with the highest expected reward $a^* \in \arg \max_k \mu_k$ with the least possible amount of trials and with probability of error $P(\hat{a} \neq a^*) \leq \delta$. For example using a decision rule, we could use m_1 where m is the machine no.1 (you can use a different variable respectively) and 1 is the amount for each time an attempt is made at pulling the lever, where $\sum_{m=1, m=2, \dots} = M$, identify M as the sum of each attempts $m_1 + m_2, \dots$ as needed, and from there you can get a ratio, sum or mean as quantitative probability and sample your formulation for each slots. You can also do $\sum_{i=1}^N - (\sum_{j=1}^n)$ where $m_1 + m_2$ equal to each a unique machine slot, x, y is the amount each time the lever is triggered, N is the sum of $(m_1 x, y) + (m_2 x, y) \dots$, k would be the total available amount in your possession, k is relative to N where $N = n(a, b), (n_1 a, b), (n_2 a, b) \dots$ reduced n_j as the sum of each gain or loss from a, b (for example, suppose you have 100\$ that is defined as n , and a would be a gain, b is equal to a loss. From there you get your

results either positive or negative to add for N with your own specific rule) and i as the maximum you are willing to spend. It is possible to express this construction using a combination of multiple algebraic formulation, as mentioned above where you can limit with T for, or in time and so on.

Multi-armed bandit -- Part 4

An important variation of the classical regret minimization problem in multi-armed bandits is best arm identification (BAI), also known as pure exploration. This problem is crucial in various applications, including clinical trials, adaptive routing, recommendation systems, and A/B testing. In BAI, the objective is to identify the arm having the highest expected reward. An algorithm in this setting is characterized by a sampling rule, a decision rule, and a stopping rule, described as follows:

Multi-armed bandit -- Part 5

Adversarial bandit Another variant of the multi-armed bandit problem is called the adversarial bandit, first introduced by Auer and Cesa-Bianchi (1998). In this variant, at each iteration, an agent chooses an arm and an adversary simultaneously chooses the payoff structure for each arm. This is one of the strongest generalizations of the bandit problem as it removes all assumptions of the distribution and a solution to the adversarial bandit problem is a generalized solution to the more specific bandit problems.

Multi-armed bandit -- Part 6

Non-stationary bandit This framework refers to the multi-armed bandit problem in a non-stationary setting (i.e., in presence of concept drift). In the non-stationary setting, it is assumed that the expected reward for an arm k can change at every time step $t \in T$: $\mu_{t-1}^k \neq \mu_t^k$. Thus, μ_t^k no longer represents the whole sequence of expected (stationary) rewards for arm k . Instead, μ^k denotes the sequence of expected rewards for arm k , defined as $\mu^k = \{\mu_t^k\}_{t=1}^T$. A dynamic oracle represents the optimal policy to be compared with other policies in the non-stationary setting. The dynamic oracle optimises the expected reward at each step $t \in T$ by always selecting the best arm, with expected reward of μ_t^* . Thus, the cumulative expected reward $D(T)$ for the dynamic oracle at final time step T is

An Analysis of Multi-armed bandit

Multi-armed bandit

defined as:

Multi-armed bandit -- Part 7

Pricing strategies Pricing strategies establish a price for each lever. For example, as illustrated with the POKER algorithm, the price can be the sum of the expected reward plus an estimation of extra future rewards that will gain through the additional knowledge. The lever of highest price is always pulled.

Multi-armed bandit -- Part 8

Collaborative bandit Approaches using multiple bandits that cooperate sharing knowledge in order to better optimize their performance started in 2013 with "A Gang of Bandits", an algorithm relying on a similarity graph between the different bandit problems to share knowledge. The need of a similarity graph was removed in 2014 by the work on the CLUB algorithm. Following this work, several other researchers created algorithms to learn multiple models at the same time under bandit feedback. For example, COFIBA was introduced by Li and Karatzoglou and Gentile (SIGIR 2016), where the classical collaborative filtering, and content-based filtering methods try to learn a static recommendation model given training data.