

An Analysis of Variational autoencoder

Variational autoencoder

$$\mathcal{L} = \mathbb{E}_{q_\phi}[\log p_\theta(x|z)] - D_{KL}(q_\phi(z|x) \| p(z))$$

Variational autoencoder -- Part 1

To efficiently search for $\theta^*, \phi^* = \operatorname{argmax}_{\theta, \phi} \mathcal{L}_{\theta, \phi}(x)$ the typical method is gradient ascent. It is straightforward to find $\nabla_{\theta} \mathbb{E}_{z \sim q_{\phi}(\cdot|x)} [\ln p_{\theta}(x, z) q_{\phi}(z|x)] = \mathbb{E}_{z \sim q_{\phi}(\cdot|x)} [\nabla_{\theta} \ln p_{\theta}(x, z) q_{\phi}(z|x)]$ However, $\nabla_{\phi} \mathbb{E}_{z \sim q_{\phi}(\cdot|x)} [\ln p_{\theta}(x, z) q_{\phi}(z|x)]$ does not allow one to put the ∇_{ϕ} inside the expectation, since ϕ appears in the probability distribution itself. The reparameterization trick (also known as stochastic backpropagation) bypasses this difficulty. The most important example is when $z \sim q_{\phi}(\cdot|x)$ is normally distributed, as $N(\mu_{\phi}(x), \Sigma_{\phi}(x))$.

Variational autoencoder -- Part 2

the sliced Wasserstein distance used by S Kolouri, et al. in their VAE the energy distance implemented in the Radon Sobolev Variational Auto-Encoder the Maximum Mean Discrepancy distance used in the MMD-VAE the Wasserstein distance used in the WAEs kernel-based distances used in the Kernelized Variational Autoencoder (K-VAE)

Variational autoencoder -- Part 3

where $p_{\theta}(x, z)$ represents the joint distribution under p_{θ} of the observable data x and its latent representation or encoding z . According to the chain rule, the equation can be rewritten as

Variational autoencoder -- Part 4

the usual reconstruction error part which seeks to ensure that the encoder-then-decoder mapping $x \mapsto D_{\theta}(E_{\psi}(x))$ is as close to

the identity map as possible; the sampling is done at run time from the empirical distribution $\mathbb{P}^{\text{real}}$ of objects available (e.g., for MNIST or IMAGENET this will be the empirical probability law of all images in the dataset). This gives the term: $\mathbb{E}_{x \sim \mathbb{P}^{\text{real}}} [D_{\theta}(E_{\phi}(x))]^2$ a variational part that ensures that, when the empirical distribution $\mathbb{P}^{\text{real}}$ is passed through the encoder E_{ϕ} , we recover the target distribution, denoted here $\mu(dz)$ that is usually taken to be a Multivariate normal distribution. We will denote $E_{\phi} \sharp \mathbb{P}^{\text{real}}$ this pushforward measure which in practice is just the empirical distribution obtained by passing all dataset objects through the encoder E_{ϕ} . In order to make sure that $E_{\phi} \sharp \mathbb{P}^{\text{real}}$ is close to the target $\mu(dz)$, a Statistical distance d is invoked and the term $d(\mu(dz), E_{\phi} \sharp \mathbb{P}^{\text{real}})^2$ is added to the loss. We obtain the final formula for the loss:

Variational autoencoder -- Part 5

The statistical distance d requires special properties, for instance it has to possess a formula as expectation because the loss function will need to be optimized by stochastic optimization algorithms. Several distances can be chosen and this gave rise to several flavors of VAEs:

Variational autoencoder -- Part 6

$$D_{KL}(q_{\phi}(z|x) \| p_{\theta}(z|x)) = \mathbb{E}_{z \sim q_{\phi}(\cdot|x)} [\ln q_{\phi}(z|x) p_{\theta}(z|x)] = \mathbb{E}_{z \sim q_{\phi}(\cdot|x)} [\ln q_{\phi}(z|x) p_{\theta}(x, z)] = \ln p_{\theta}(x) + \mathbb{E}_{z \sim q_{\phi}(\cdot|x)} [\ln q_{\phi}(z|x) p_{\theta}(x, z)] \\ \begin{aligned} D_{KL}(q_{\phi}(\cdot|x) \| p_{\theta}(\cdot|x)) &= \mathbb{E}_{z \sim q_{\phi}(\cdot|x)} [\ln q_{\phi}(z|x) p_{\theta}(x, z)] - \mathbb{E}_{z \sim q_{\phi}(\cdot|x)} [\ln p_{\theta}(x, z)] \\ &= \mathbb{E}_{z \sim q_{\phi}(\cdot|x)} [\ln q_{\phi}(z|x) p_{\theta}(x, z)] - \ln p_{\theta}(x) \\ &= \mathbb{E}_{z \sim q_{\phi}(\cdot|x)} [\ln q_{\phi}(z|x) p_{\theta}(x, z)] - \ln p_{\theta}(x) \end{aligned}$$

Variational autoencoder -- Part 7

Now, define the evidence lower bound (ELBO): $\mathcal{L}_{\theta, \phi}(x) := \mathbb{E}_{z \sim q_{\phi}(\cdot|x)} [\ln p_{\theta}(x, z) q_{\phi}(z|x)] = \ln p_{\theta}(x) - D_{KL}(q_{\phi}(\cdot|x) \| p_{\theta}(\cdot|x))$

An Analysis of Variational autoencoder

Variational autoencoder

$\mathbb{E}_{z \sim q_{\phi}(\cdot|x)} \left[\ln \frac{p_{\theta}(x,z)}{q_{\phi}(z|x)} \right] = \ln p_{\theta}(x) - D_{KL}(q_{\phi}(\cdot|x) \parallel p_{\theta}(\cdot|x))$ Maximizing the ELBO $\theta^*, \phi^* = \operatorname{argmax}_{\theta, \phi} L_{\theta, \phi}(x)$ is equivalent to simultaneously maximizing $\ln p_{\theta}(x)$ and minimizing $D_{KL}(q_{\phi}(z|x) \parallel p_{\theta}(z|x))$. That is, maximizing the log-likelihood of the observed data, and minimizing the divergence from the approximate posterior $q_{\phi}(\cdot|x)$ to the exact posterior $p_{\theta}(\cdot|x)$. The form given is not very convenient for maximization, but the following, equivalent form, is: $L_{\theta, \phi}(x) = \mathbb{E}_{z \sim q_{\phi}(\cdot|x)} [\ln p_{\theta}(x|z)] - D_{KL}(q_{\phi}(\cdot|x) \parallel p_{\theta}(\cdot|x))$ where $\ln p_{\theta}(x|z)$ is implemented as $-\frac{1}{2} \|x - D_{\theta}(z)\|_2^2$, since that is, up to an additive constant, what $x|z \sim N(D_{\theta}(z), I)$ yields. That is, we model the distribution of x conditional on z to be a Gaussian distribution centered on $D_{\theta}(z)$. The distribution of $q_{\phi}(z|x)$ and $p_{\theta}(z)$ are often also chosen to be Gaussians as $z|x \sim N(E_{\phi}(x), \sigma_{\phi}(x)^2 I)$ and $z \sim N(0, I)$, with which we obtain by the formula for KL divergence of Gaussians: $L_{\theta, \phi}(x) = -\frac{1}{2} \mathbb{E}_{z \sim q_{\phi}(\cdot|x)} [\|x - D_{\theta}(z)\|_2^2] - \frac{1}{2} (N \sigma_{\phi}(x)^2 + \|E_{\phi}(x)\|_2^2 - 2N \ln \sigma_{\phi}(x)) + \text{const}$ Here N is the dimension of z . For a more detailed derivation and more interpretations of ELBO and its maximization, see its main page.