

Foundations of Lasso regression

[Lasso regression]

$$L(\beta) = \|y - X\beta\|^2 + \lambda \|\beta\|_1$$

1.0 Content Block

where q_i is specified below and the "prime" symbol stands for transpose. The first fraction represents relative accuracy, the second fraction relative simplicity, and λ balances between the two.

2.0 Content Block

where $y \sim 0 = y - X\beta_0$. If $b = b_{OLS}$ when R^2 is computed, then the diagonal elements of $R \otimes$ sum to R^2 . The diagonal $R \otimes$ values may be smaller than 0 or, less often, larger than 1. If regressors are uncorrelated, then the i^{th} diagonal element of $R \otimes$ simply corresponds to the R^2 value between x_i and y . A rescaled version of the adaptive lasso of can be obtained by setting $q_{\text{adaptive lasso}, i} = |b_{OLS, i} - \beta_0|$. If regressors are uncorrelated, the moment that the i^{th} parameter is activated is given by the i^{th} diagonal element of $R \otimes$. Assuming for convenience that β_0 is a vector of zeros,

3.0 Content Block

$\min_{\beta \in \mathbb{R}^p} \left\{ \frac{1}{N} \|y - X\beta\|^2 + \lambda \|\beta\|_1 \right\}$, $\left\{ \frac{1}{N} \|y - X\beta\|^2 + \lambda \|\beta\|_1 \right\}$

4.0 Content Block

$\min_{\beta \in \mathbb{R}^p} \left\{ \frac{1}{N} \|y - X\beta\|^2 + \lambda \sum_{j=1}^p \vartheta(\beta_j) \right\}$, $\left\{ \frac{1}{N} \|y - X\beta\|^2 + \lambda \sum_{j=1}^p \vartheta(\beta_j) \right\}$

5.0 Content Block

That is, if regressors are uncorrelated, λ again specifies the minimal influence of β_0 . Even when regressors are correlated, the first time that a regression parameter is activated occurs when λ is equal to the highest diagonal element of $R \otimes$. These results can be compared to a rescaled version of the lasso by defining $q_{\text{lasso}, i} = \frac{1}{p} \sum_{j=1}^p |b_{OLS, j} - \beta_0|$, which is the average absolute deviation of b_{OLS} from β_0 . Assuming that regressors are uncorrelated, then the moment of activation of the i^{th} regressor is given by

6.0 Content Block

In statistics and machine learning, lasso (least absolute shrinkage and selection operator; also Lasso, LASSO or L1 regularization) is a regression analysis method that performs both variable selection and regularization in order to enhance the prediction accuracy and interpretability of the resulting statistical model. The lasso method assumes that the coefficients of the linear model are sparse, meaning that few of them are non-zero. It was originally introduced in geophysics, and later by Robert Tibshirani, who coined the term. Lasso was originally formulated for linear regression models. This simple case reveals a substantial amount about the estimator. These include its relationship to ridge regression and best subset selection and the connections between lasso coefficient estimates and so-called soft thresholding. It also reveals that (like standard linear regression) the coefficient estimates do not need to be unique if covariates are collinear. Though originally defined for linear regression, lasso regularization is easily extended to other statistical models including generalized linear models, generalized estimating equations, proportional hazards models, and M-estimators. Lasso's ability to perform subset selection relies on the form of the constraint and has a variety of interpretations including in terms of geometry, Bayesian statistics and convex analysis. The LASSO is closely related to basis pursuit denoising.

7.0 Content Block

where $\vartheta(\gamma)$ is an arbitrary concave monotonically increasing function (for example, $\vartheta(\gamma) = \sqrt{\gamma}$ gives the lasso penalty and $\vartheta(\gamma) = \gamma^{1/4}$ gives the $\ell_{1/2}$ penalty). The efficient algorithm for minimization is based on piece-wise quadratic approximation of subquadratic growth (PQSQ).

8.0 Content Block

$\min_{\beta} \|\beta\|_1$ subject to $(1 - \alpha) \|\beta\|_1 + \alpha \|\beta\|_2^2 \leq t$, where $\alpha = \frac{\lambda_2}{\lambda_1 + \lambda_2}$.

9.0 Content Block

Lasso can set coefficients to zero, while the superficially similar ridge regression cannot. This is due to the difference in the shape of their constraint boundaries. Both lasso and ridge regression can be interpreted as minimizing the same objective function

10.0 Content Block

For $p = 1$, the moment of activation is again given by $\lambda \sim \text{lasso}$, $i = R^2$. If β_0 is a vector of zeros and a subset of p_B relevant parameters are equally responsible for a perfect fit of $R^2 = 1$, then this subset is activated at a λ value of $\frac{1}{p}$. The moment of activation of a relevant regressor then equals $\frac{1}{p_B}$. In other words, the inclusion of irrelevant regressors delays the moment that relevant regressors are activated by this rescaled lasso. The adaptive lasso and the lasso are special cases of a '1ASTc' estimator. The latter only groups parameters together if the absolute correlation among regressors is larger than a user-specified value.

11.0 Content Block

Using $X^T X = I$ and the ridge regression formula: $\hat{\beta} = (X^T X + N \lambda I)^{-1} X^T y$, yields: