

Research Notes on Diffusion model

Diffusion model

$$q(x_t | x_{t-1}) = \mathcal{N}(\sqrt{1 - \beta_t} x_{t-1}, \beta_t I)$$

>> Section 1

$d \phi_t(x) = v_t(\phi_t(x))$ $t \in [0, 1]$, starting from $\phi_0(x) = x$ $\frac{d}{dt} \phi_t(x) = v_t(\phi_t(x))$ $t \in [0, 1]$, $\phi_0(x) = x$

>> Section 2

Further reading Review papers Yang, Ling (2024-09-06), YangLing0818/Diffusion-Models-Papers-Survey-Taxonomy, retrieved 2024-09-06 Yang, Ling; Zhang, Zhilong; Song, Yang; Hong, Shenda; Xu, Runsheng; Zhao, Yue; Zhang, Wentao; Cui, Bin; Yang, Ming-Hsuan (2023-11-09). "Diffusion Models: A Comprehensive Survey of Methods and Applications". ACM Comput. Surv. 56 (4): 105:1–105:39. arXiv:2209.00796. doi:10.1145/3626235. ISSN 0360-0300. Austin, Jacob; Johnson, Daniel D.; Ho, Jonathan; Tarlow, Daniel; Rianne van den Berg (2021). "Structured Denoising Diffusion Models in Discrete State-Spaces". arXiv:2107.03006 [cs.LG]. Croitoru, Florinel-Alin; Hondru, Vlad; Ionescu, Radu Tudor; Shah, Mubarak (2023-09-01). "Diffusion Models in Vision: A Survey". IEEE Transactions on Pattern Analysis and Machine Intelligence. 45 (9): 10850–10869. arXiv:2209.04747. Bibcode:2023ITPAM..4510850C. doi:10.1109/TPAMI.2023.3261988. ISSN 0162-8828. PMID 37030794. Mathematical details omitted in the article. "Power of Diffusion Models". AstraBlog. 2022-09-25. Retrieved 2023-09-25. Luo, Calvin (2022-08-25). "Understanding Diffusion Models: A Unified Perspective". arXiv:2208.11970 [cs.LG]. Weng, Lilian (2021-07-11). "What are Diffusion Models?". lilianweng.github.io. Retrieved 2023-09-25. Tutorials Nakkiran, Preetum; Bradley, Arwen; Zhou, Hattie; Advani, Madhu (2024). "Step-by-Step Diffusion: An Elementary Tutorial". arXiv:2406.08929 [cs.LG]. "Guidance: a cheat code for diffusion models". 26 May 2022. Overview of classifier guidance and classifier-free guidance, light on mathematical details. Catherine Higham, Desmond J. Higham, and Peter Grindrod: "Diffusion Models for Generative Artificial Intelligence: An Introduction for Applied Mathematicians", SIAM Review, Vol.67, No.3 (2025).

>> Section 3

$L_{\text{simple}, t} = E_{x_0 \sim q; z \sim N(0, I)} [\|\theta(x_t, t) - z\|_2^2]$ $L_{\text{simple}, t} = E_{x_0 \sim q; z \sim N(0, I)} \left[\left\| \theta(x_t, t) - z \right\|_2^2 \right]$

$$\|x_t - z\|_2^2$$

>> Section 4

$p_t(x) = \int p_t(x|z) q(z) dz$ and $v_t(x) = E_{q(z)} [v_t(x|z)]$ $p_t(x) = \int p_t(x|z) q(z) dz$ and $v_t(x) = E_{q(z)} [v_t(x|z)]$

>> Section 5

In particular, we see that we can directly sample from any point in the continuous diffusion process without going through the intermediate steps, by first sampling $x_0 \sim q, z \sim N(0, I)$ $x_0 \sim q, z \sim N(0, I)$, then get $x_t = e^{-\frac{1}{2} \int_0^t \beta(s) ds} x_0 + (1 - e^{-\frac{1}{2} \int_0^t \beta(s) ds}) z$ $x_t = e^{-\frac{1}{2} \int_0^t \beta(s) ds} x_0 + (1 - e^{-\frac{1}{2} \int_0^t \beta(s) ds}) z$. That is, we can quickly sample $x_t \sim p_t$ $x_t \sim p_t$ for any $t \geq 0$ $t \geq 0$. Now, define a certain probability distribution γ γ over $[0, \infty)$ $[0, \infty)$, then the score-matching loss function is defined as the expected Fisher divergence:

>> Section 6

$L(\theta) = E_{t \sim \gamma, x_t \sim p_t} [\|\nabla f_\theta(x_t, t)\|_2^2 + 2 \nabla f_\theta(x_t, t) \cdot \nabla \log p_t(x_t, t)]$ $L(\theta) = E_{t \sim \gamma, x_t \sim p_t} [\|\nabla f_\theta(x_t, t)\|_2^2 + 2 \nabla f_\theta(x_t, t) \cdot \nabla \log p_t(x_t, t)]$

>> Section 7

$dy_t = \frac{1}{2} \beta(T-t) y_t dt + \beta(T-t) \nabla y_t \ln p_{T-t}(y_t) dt + \beta(T-t) dW_t$ $dy_t = \frac{1}{2} \beta(T-t) y_t dt + \beta(T-t) \nabla y_t \ln p_{T-t}(y_t) dt + \beta(T-t) dW_t$

>> Section 8

Noise prediction network Since $x_{t-1} | x_t, x_0 \sim N(\mu_{t-1}(x_t, x_0), \sigma_{t-1}^2)$ $x_{t-1} | x_t, x_0 \sim N(\mu_{t-1}(x_t, x_0), \sigma_{t-1}^2)$, this suggests that we should use $\mu_\theta(x_t, t) = \mu_{t-1}(x_t, x_0)$ $\mu_\theta(x_t, t) = \mu_{t-1}(x_t, x_0)$; however, the network does not have access to x_0 x_0 , and so it has to estimate it instead. Now, since $x_t | x_0 \sim N(\alpha^{-1} x_0, \sigma_t^2)$ $x_t | x_0 \sim N(\alpha^{-1} x_0, \sigma_t^2)$, we may write $x_t = \alpha^{-1} x_0 + \sigma_t z$ $x_t = \alpha^{-1} x_0 + \sigma_t z$, where z z is some unknown Gaussian noise. Now we see that estimating x_0 x_0 is equivalent to estimating z z . Therefore, let the network

Research Notes on Diffusion model

Diffusion model

output a noise vector $\epsilon_{\theta}(x_t, t)$, and let it predict $\mu_{\theta}(x_t, t) = \mu_{\tilde{t}}(x_t, x_t - \sigma_t \epsilon_{\theta}(x_t, t) \alpha_t^{-1}) = x_t - \epsilon_{\theta}(x_t, t) \beta_t / \sigma_t \alpha_t$. It remains to design $\Sigma_{\theta}(x_t, t)$. The DDPM paper suggested not learning it (since it resulted in "unstable training and poorer sample quality"), but fixing it at some value $\Sigma_{\theta}(x_t, t) = \zeta_t^2 I$, where either $\zeta_t^2 = \beta_t$ or $\zeta_t^2 = \sigma_{\tilde{t}}^2$ yielded similar performance. With this, the loss simplifies to $L_t = \beta_t^2 \alpha_t \sigma_t^2 \zeta_t^2 E_{x_0 \sim q; z \sim N(0, I)} [\|\epsilon_{\theta}(x_t, t) - z\|^2] + C$ which may be minimized by stochastic gradient descent. The paper noted empirically that an even simpler loss function $L_{\text{simple}, t} = E_{x_0 \sim q; z \sim N(0, I)} [\|\epsilon_{\theta}(x_t, t) - z\|^2]$ resulted in better models.